



Generative Adversarial Networks

Eklavya Sarkar, Biometrics Security and Privacy

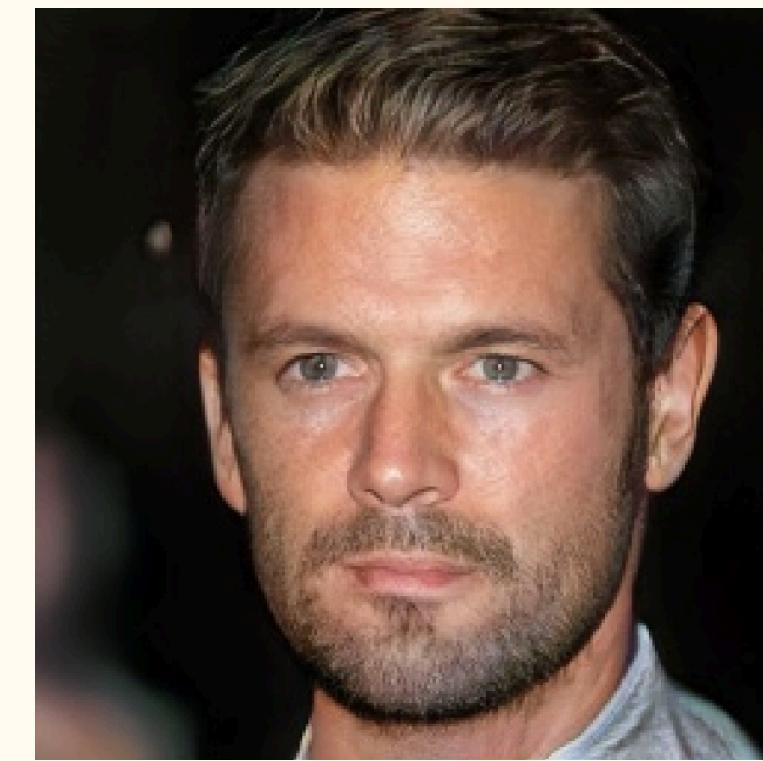
Generative Adversarial Networks (GANs)



2014



2015



2016



2017



2018



2019

Supervised Learning

1
2
3

Training Data

...
...
...

Labels

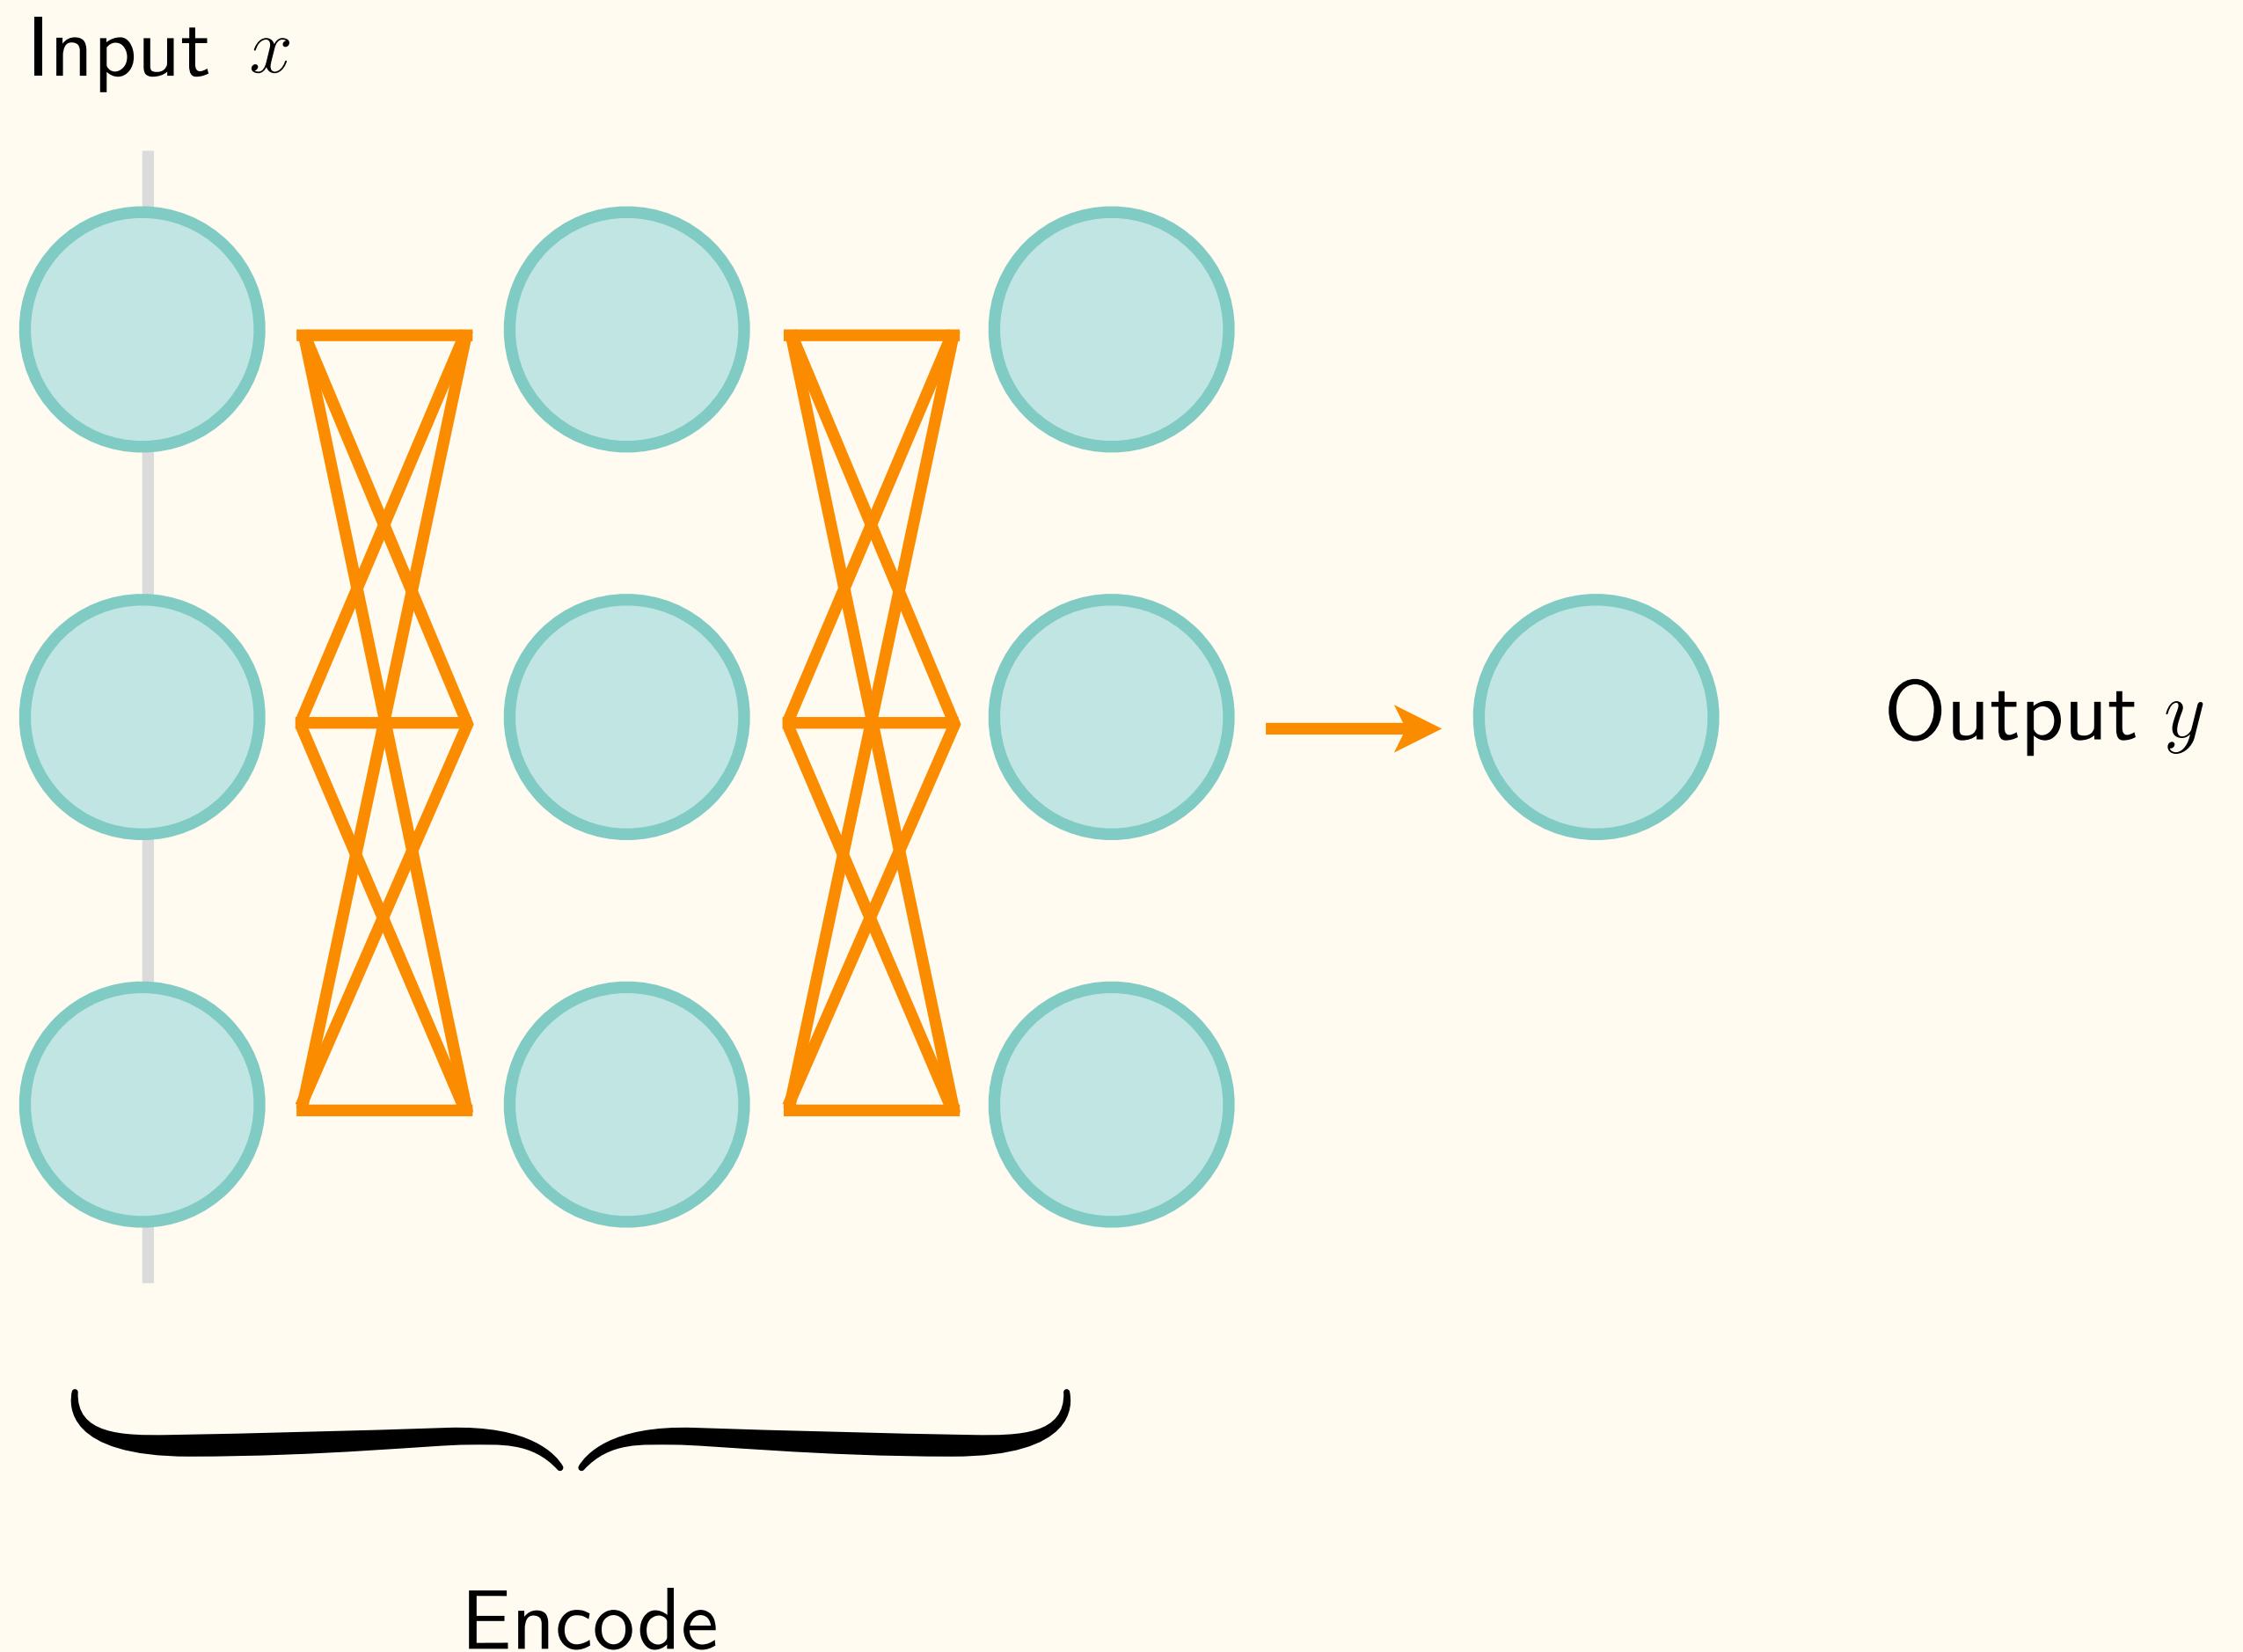
...
...
...

Predictions

Supervised Artificial
Neural Networks (ANNs):

- Take an *input*
- Produce an *output*
- Optimized with a *loss-function*

Supervised Learning



Supervised Artificial
Neural Networks (ANNs):

- Take an *input*
- Produce an *output*
- Optimized with a *loss-function*

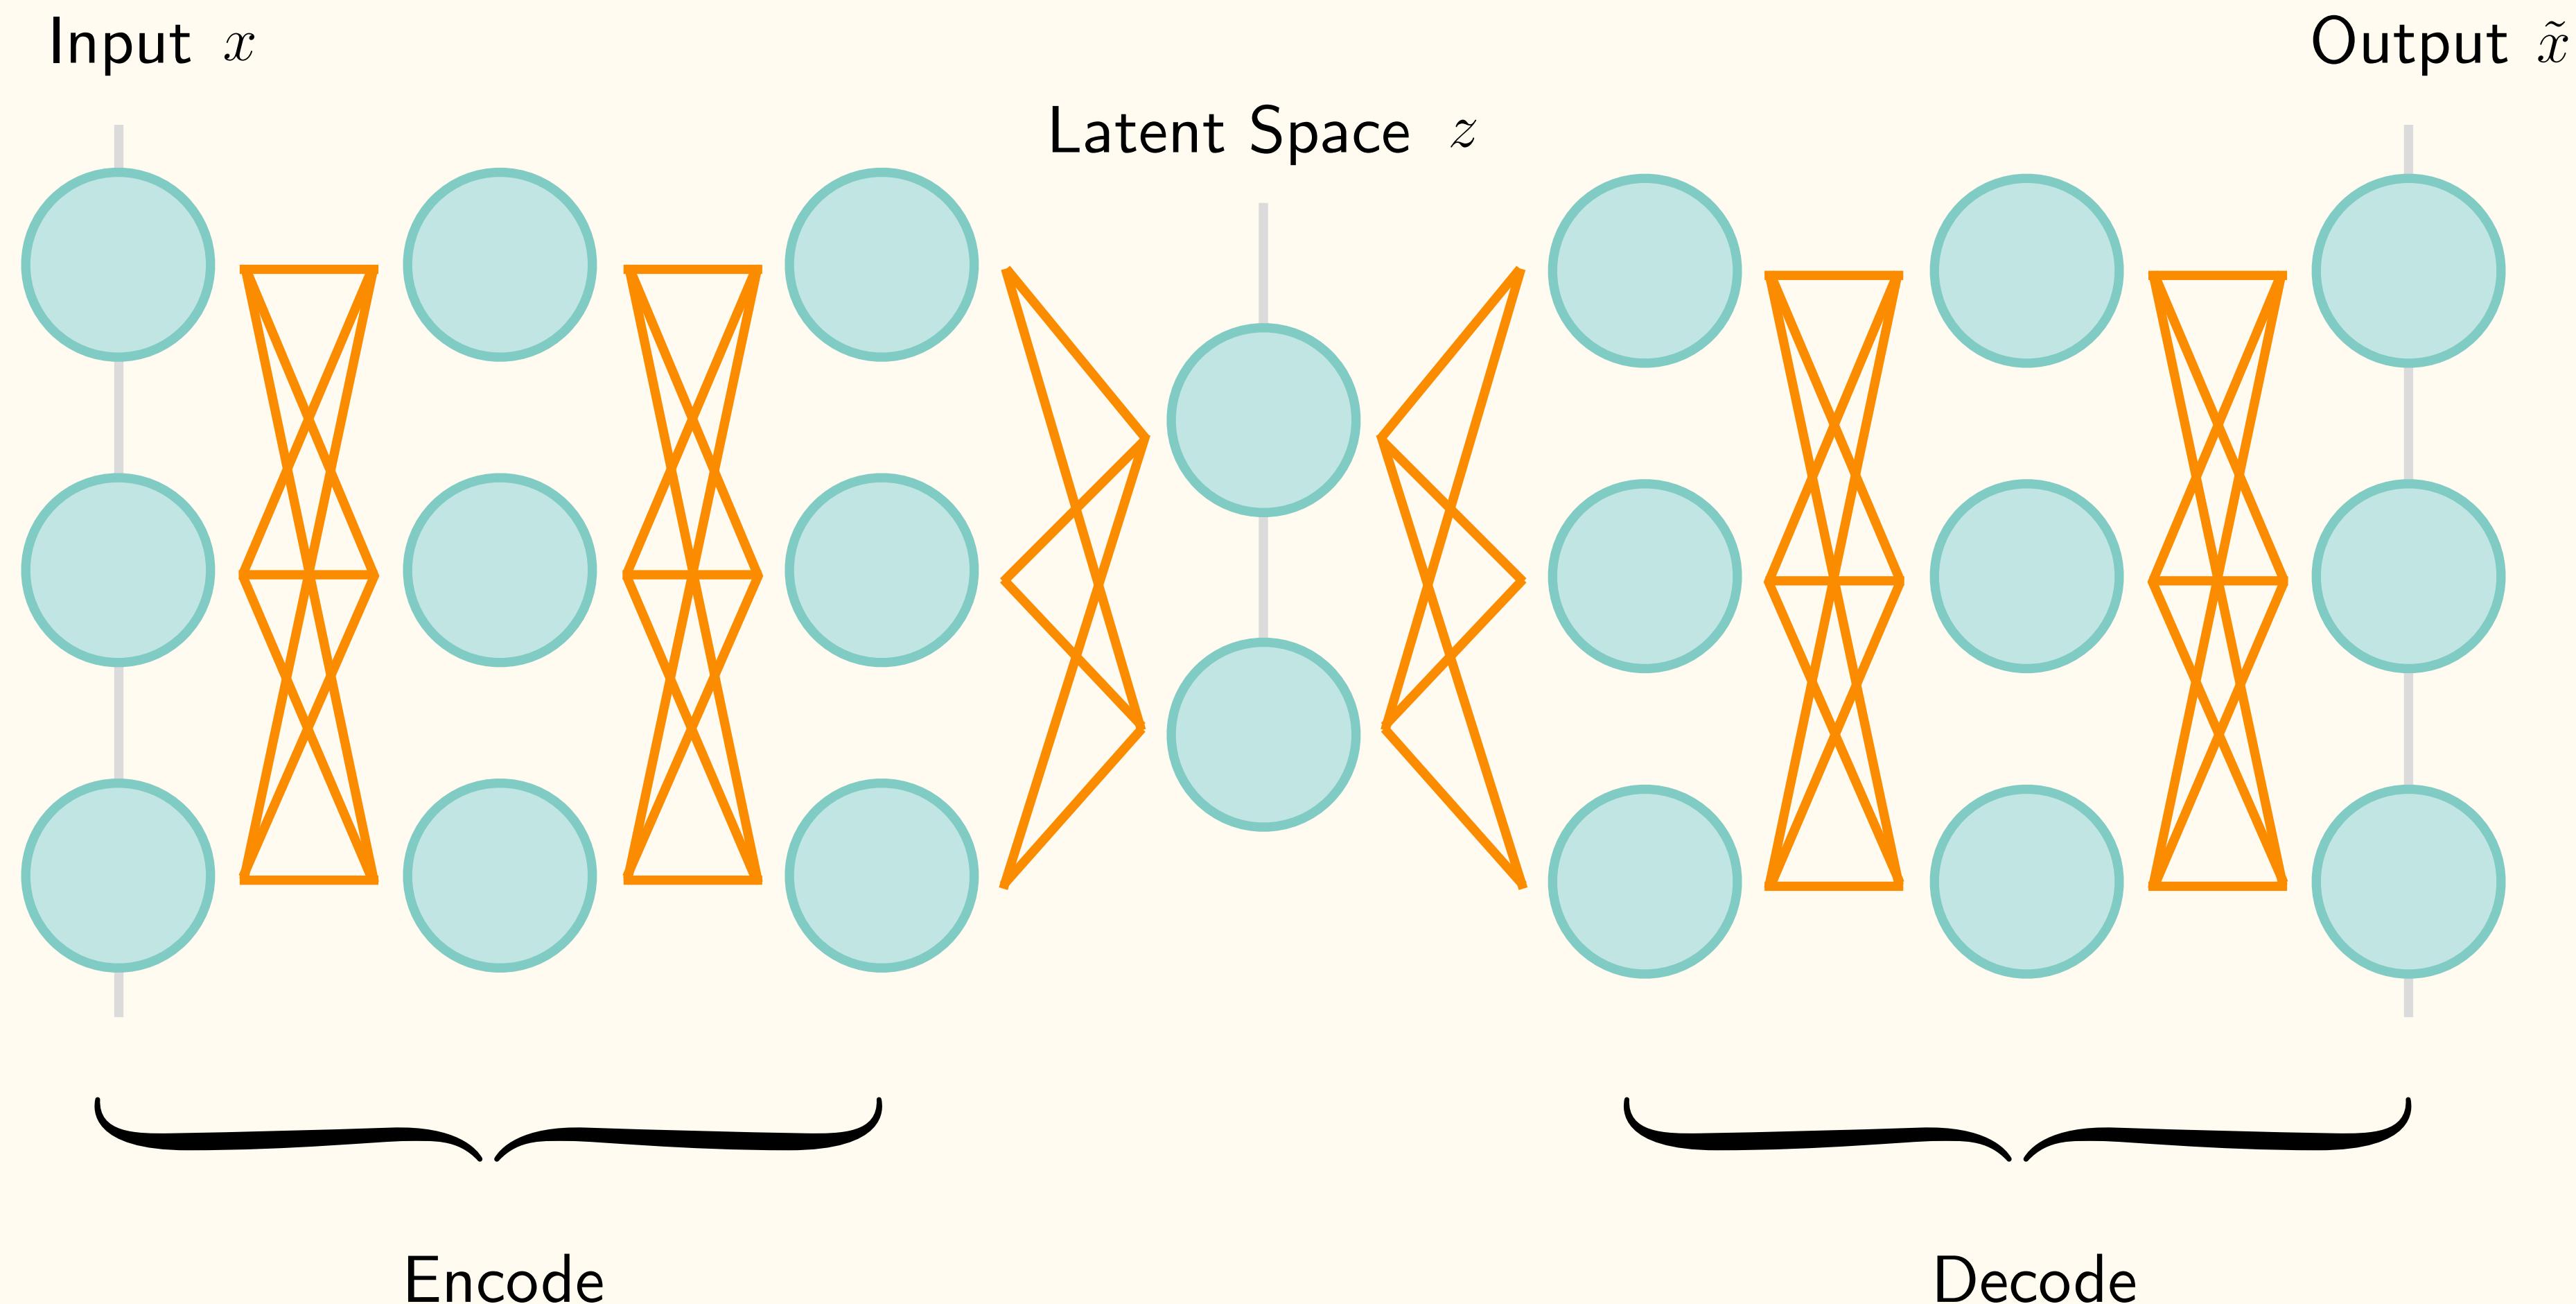
Unsupervised Learning

- Generative models are different.
- Instead of taking a sample as *input*, they produce a sample as *output*.

Unsupervised Learning

- Train model on photographs of cats
 - ▶ Learn to produce new cat-like images
- Train it on known drug molecules
 - ▶ Generate new ‘drug-like’ molecules to use as candidates in a virtual screen.
- A generative model is trained on samples which are drawn from a:
probability distribution
- Its job is to produce new samples from that same probability distribution.

Auto-Encoder



Auto-Encoder

- **Encoder:** samples → compressed representations
- **Decoder:** compressed representations in latent space → original samples
- What if we take *random vectors* in the latent space (picking a random value for each component of the vector) and pass them through the decoder ?
- ▶ If everything goes well, the decoder should produce a completely *new sample* that *resembles* the ones it was trained on.

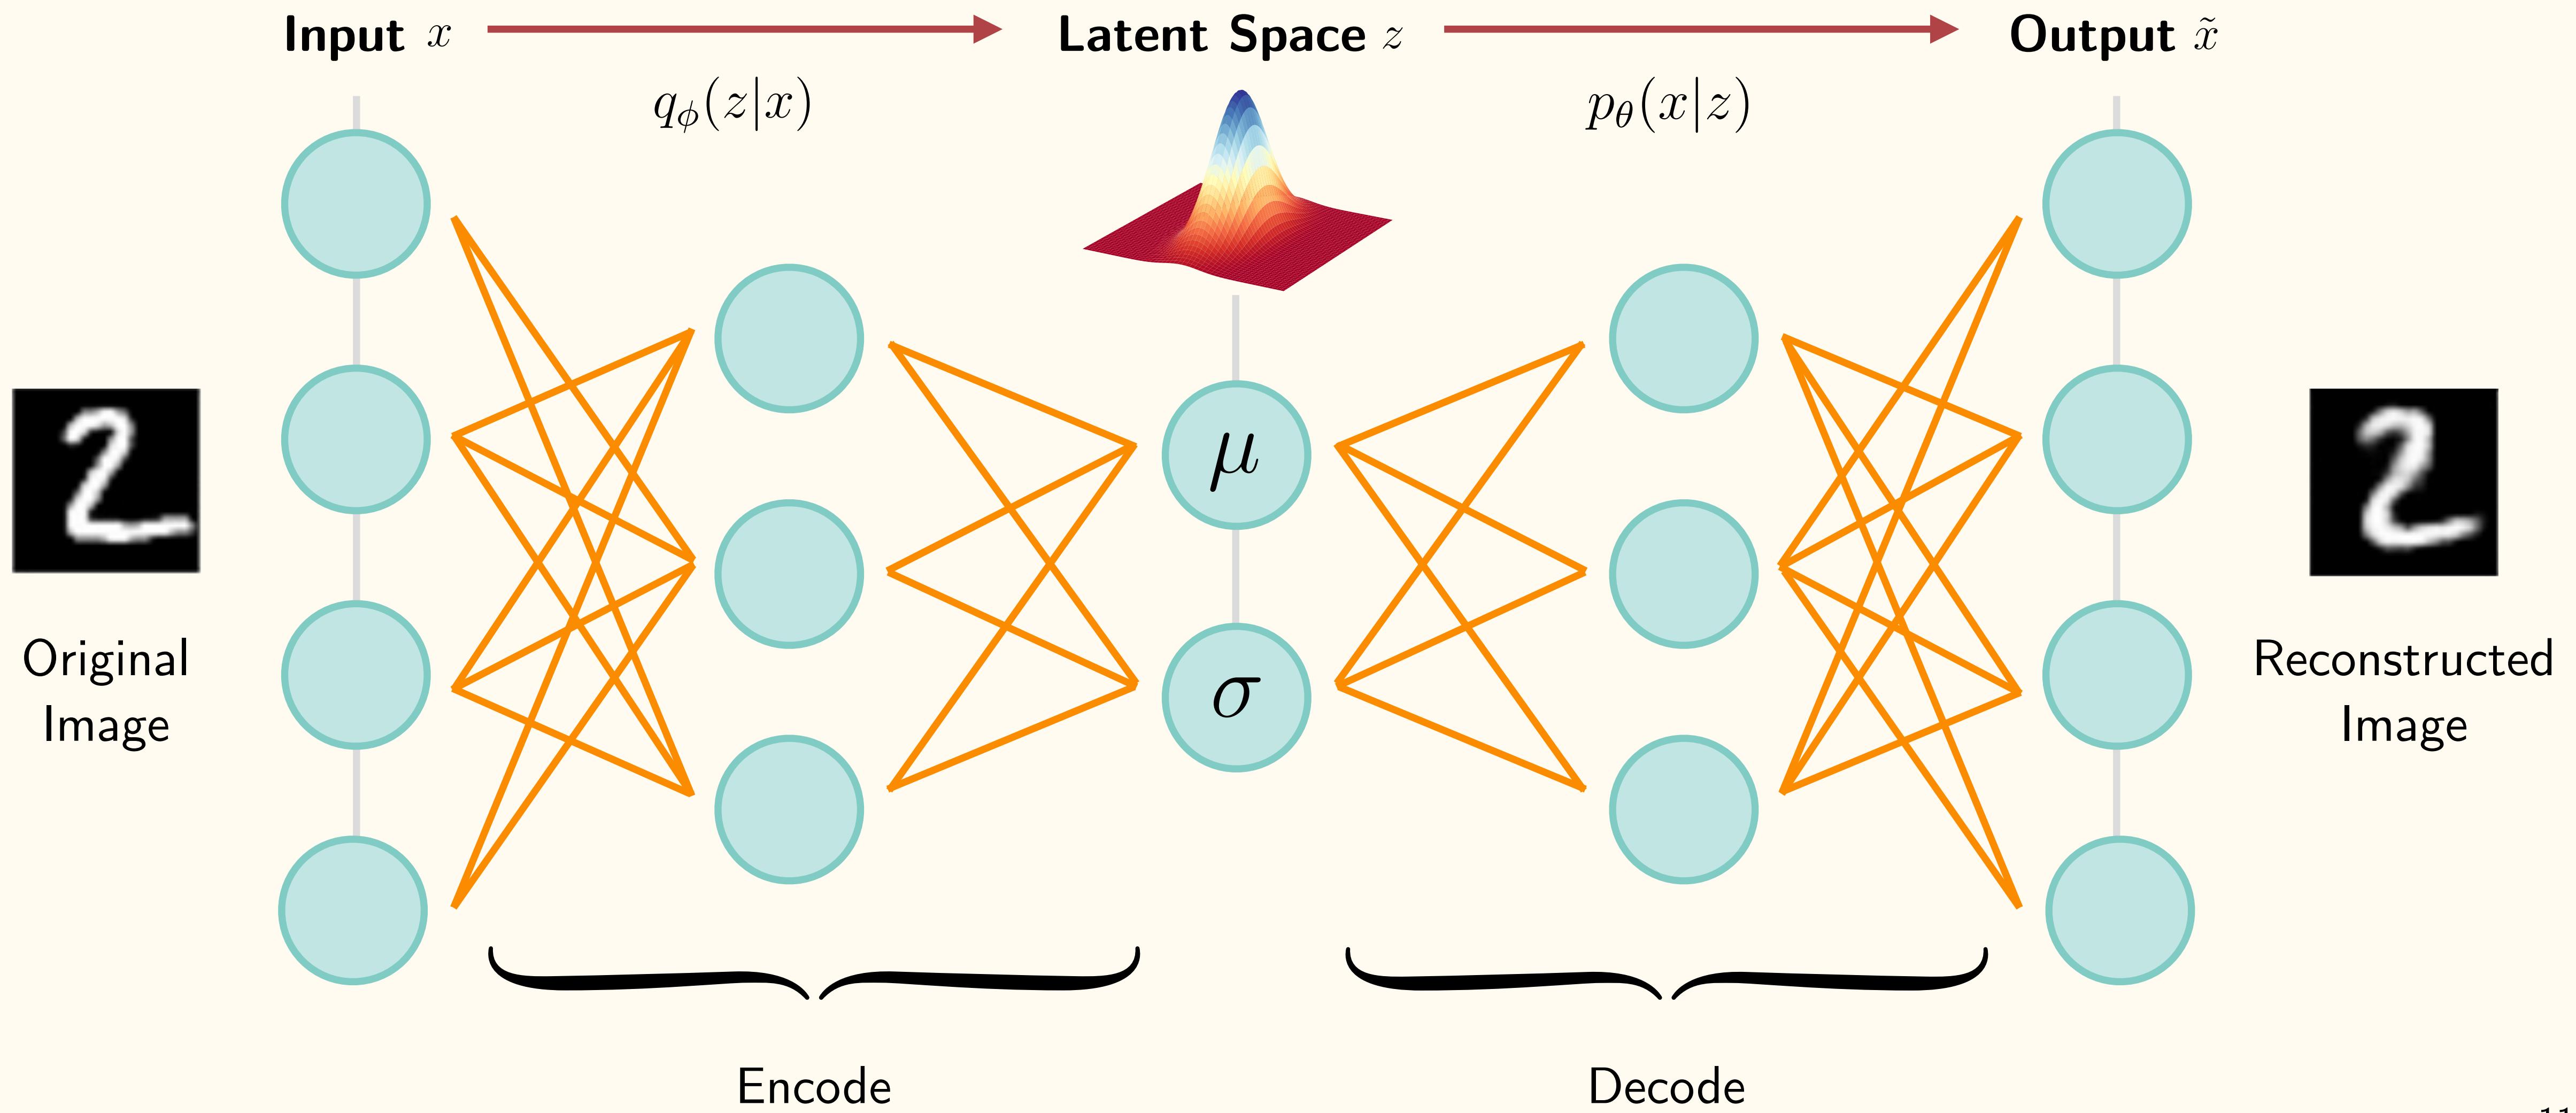
Auto-Encoder

- Problem ?
 - ▶ Doesn't work very well.
- Encoder may only produce vectors in a *small region* of the latent space.
- But if we pick a vector from *outside* that region ?
 - ▶ Output could look nothing like the training samples
- Decoder has only learned to work for the *particular* latent vectors from encoder
 - ▶ Not for *arbitrary* ones.

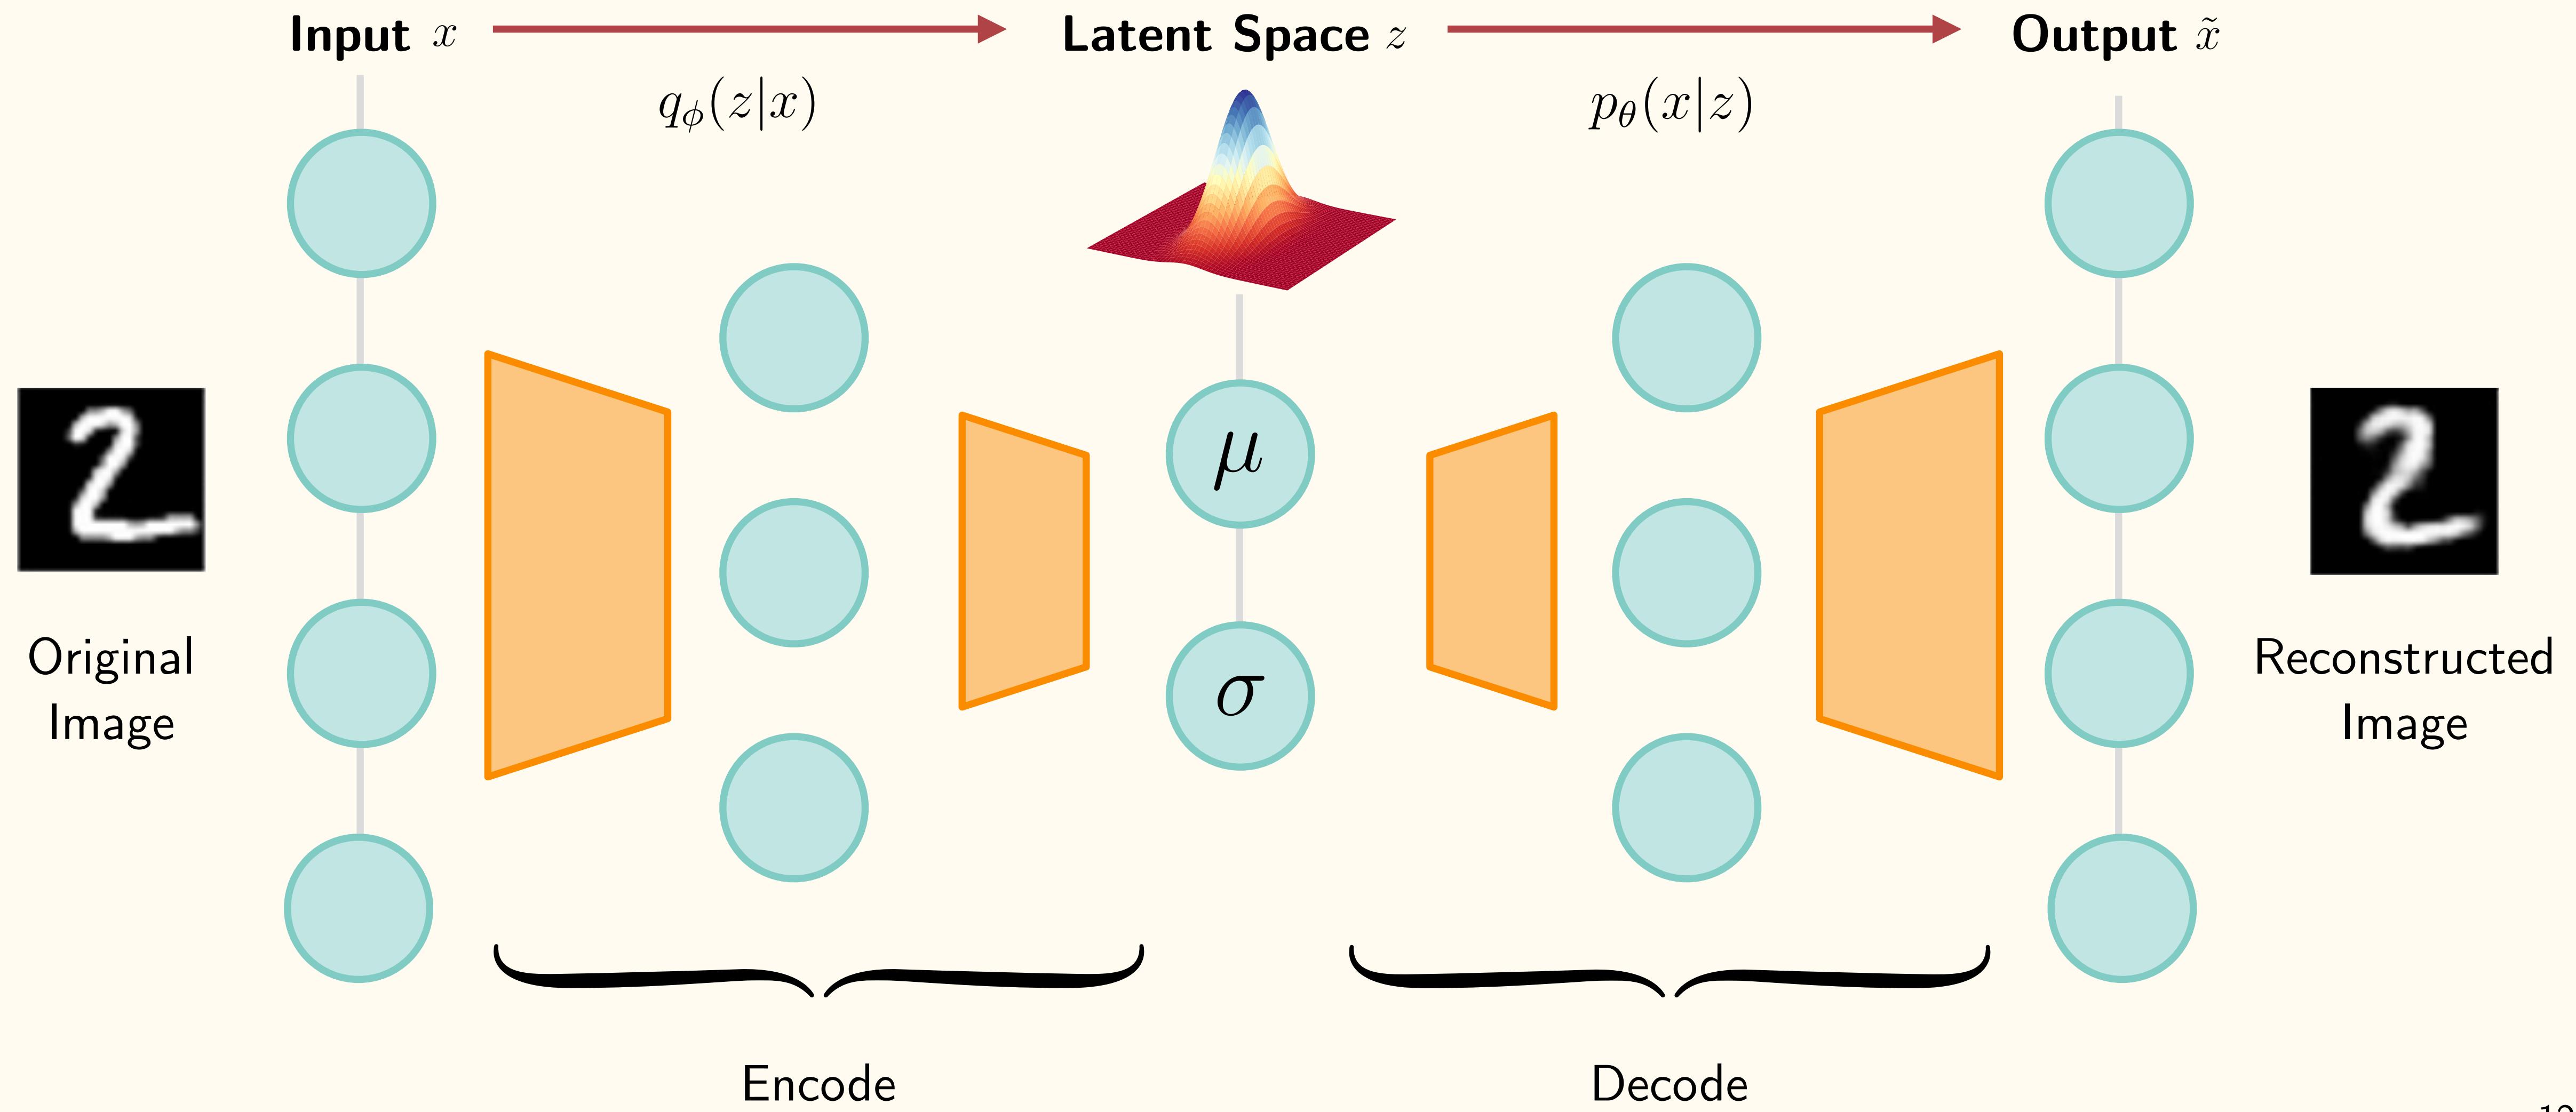
Variational Auto-Encoder (VAE)

- Solution ?
 - ▶ Variational Auto-Encoder
- Add a term to the loss function that *forces* the latent vectors to **follow a specified distribution**, e.g. $\mathcal{N} \sim (0,1)$
 - ▶ We can *expect* the decoder to work well on them
- We add random noise to the latent vector
 - ▶ Prevents decoder from being too sensitive to details

Variational Auto-Encoder (VAE)

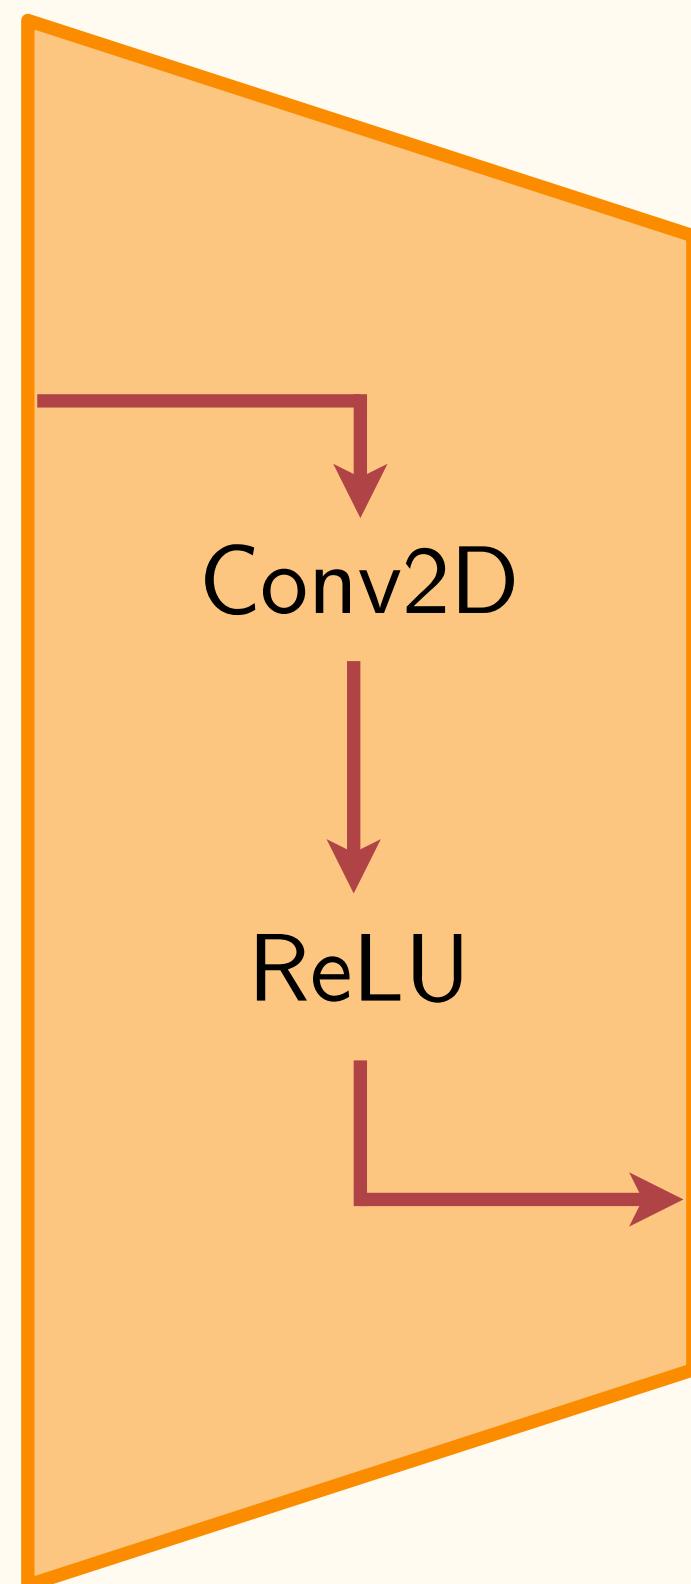


Simplified VAE Representation

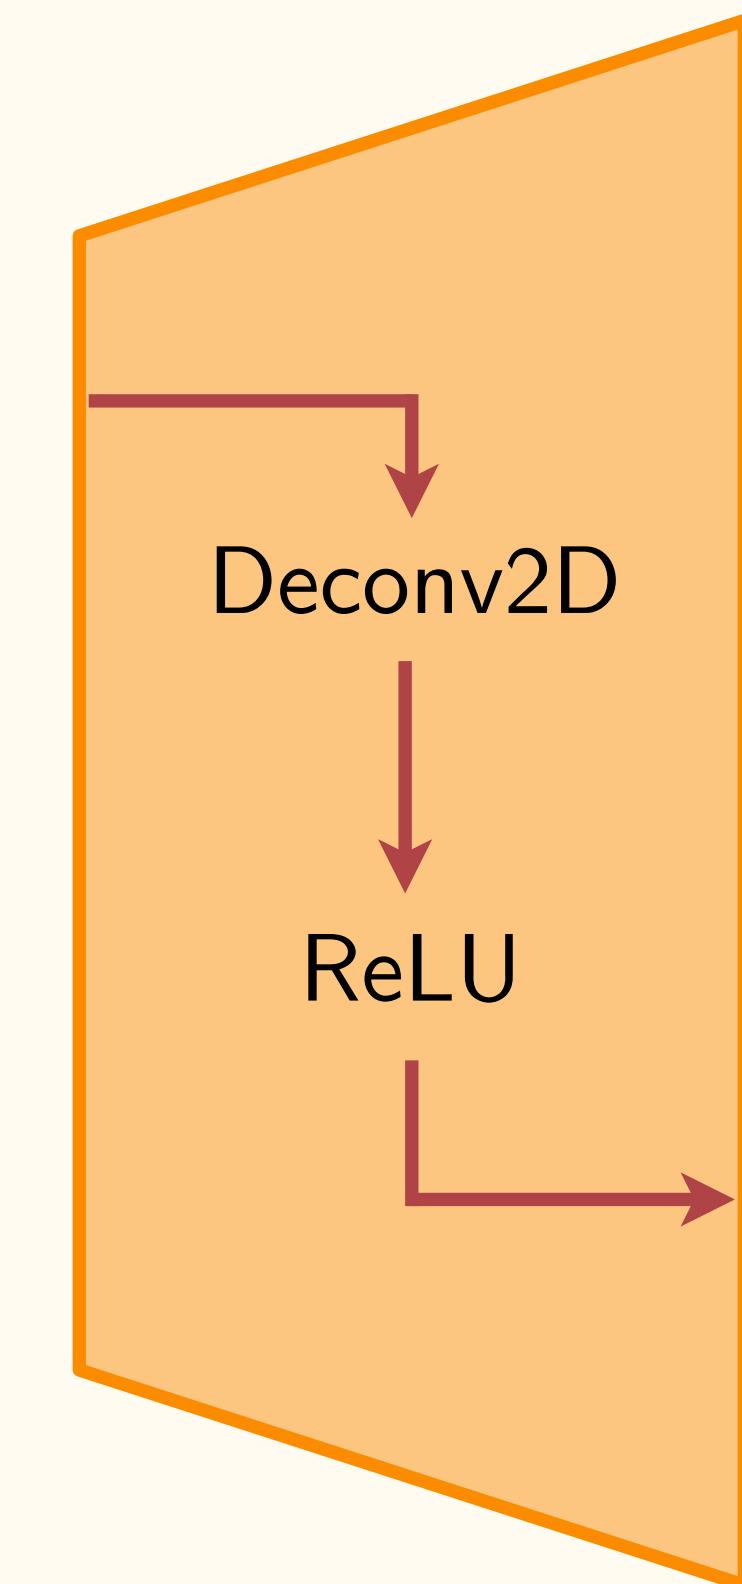


Encoder-Decoder

Encode



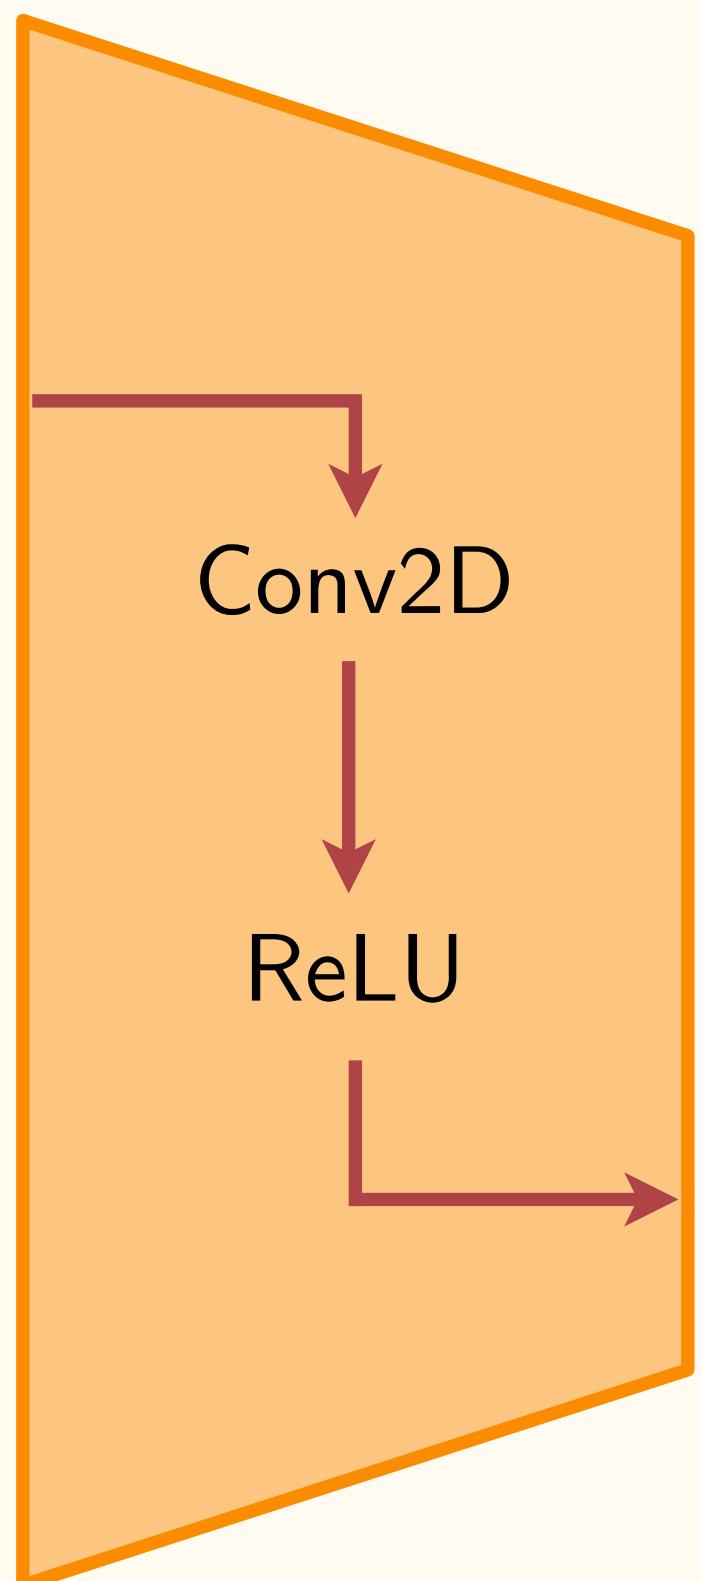
Decode



- An **encoder** is just a convolutional operation coupled with an activation function
- A **decoder** is de-convolutional operation and activation function
- De-conv2D is just a **transposed** Conv2D

Encoder-Decoder

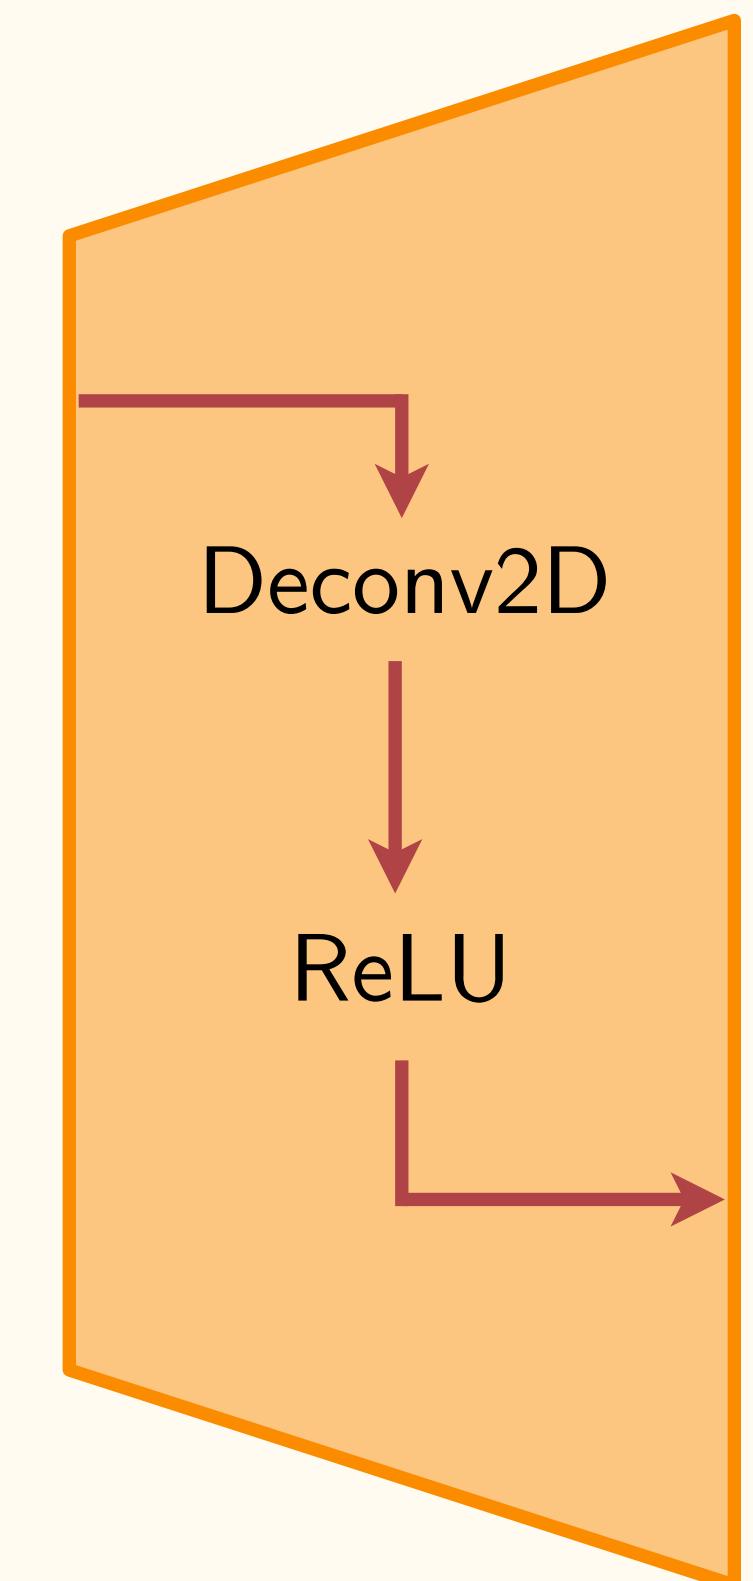
Encode



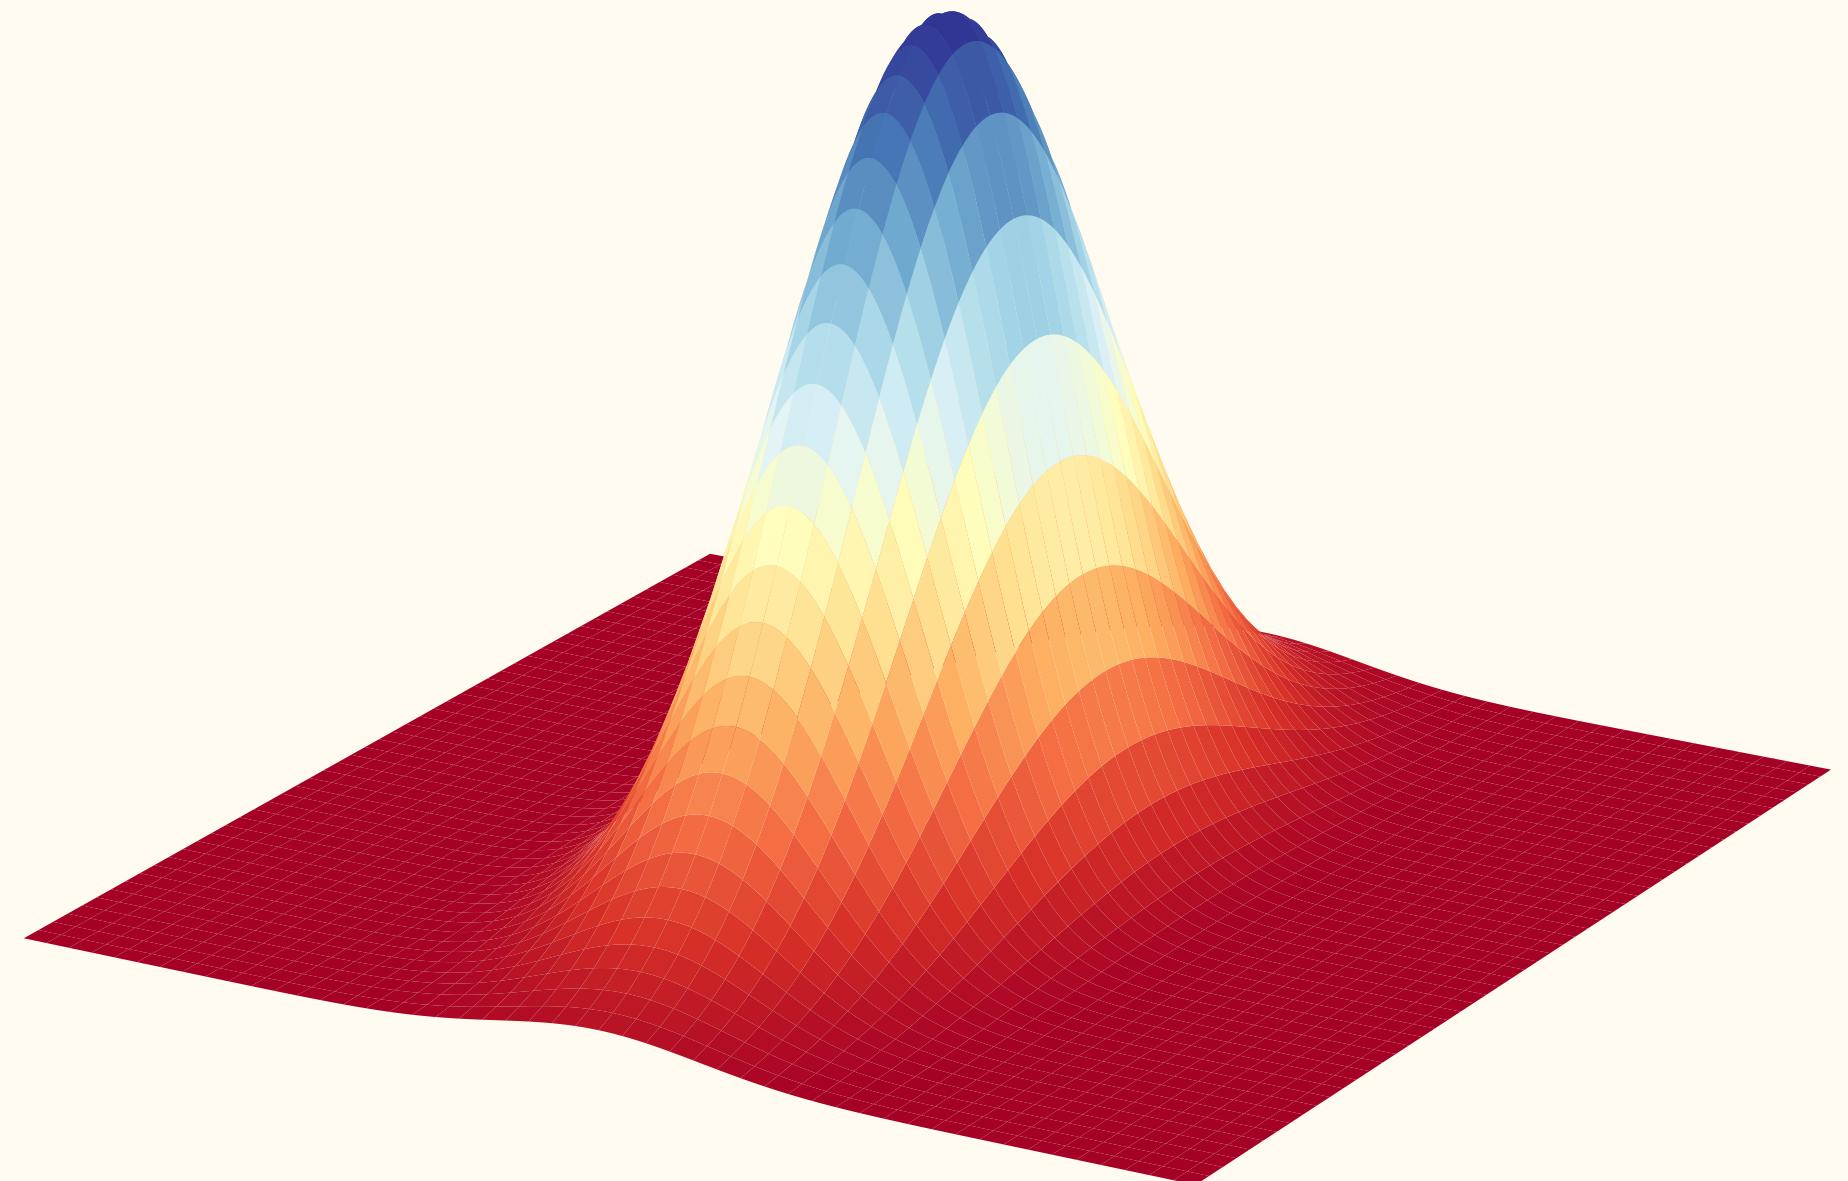
```
# Encoder  
tf.nn.conv2d()  
tf.nn.relu()
```

```
# Decoder  
tf.nn.conv2d_transpose()  
tf.nn.relu()
```

Decode



Latent Distribution

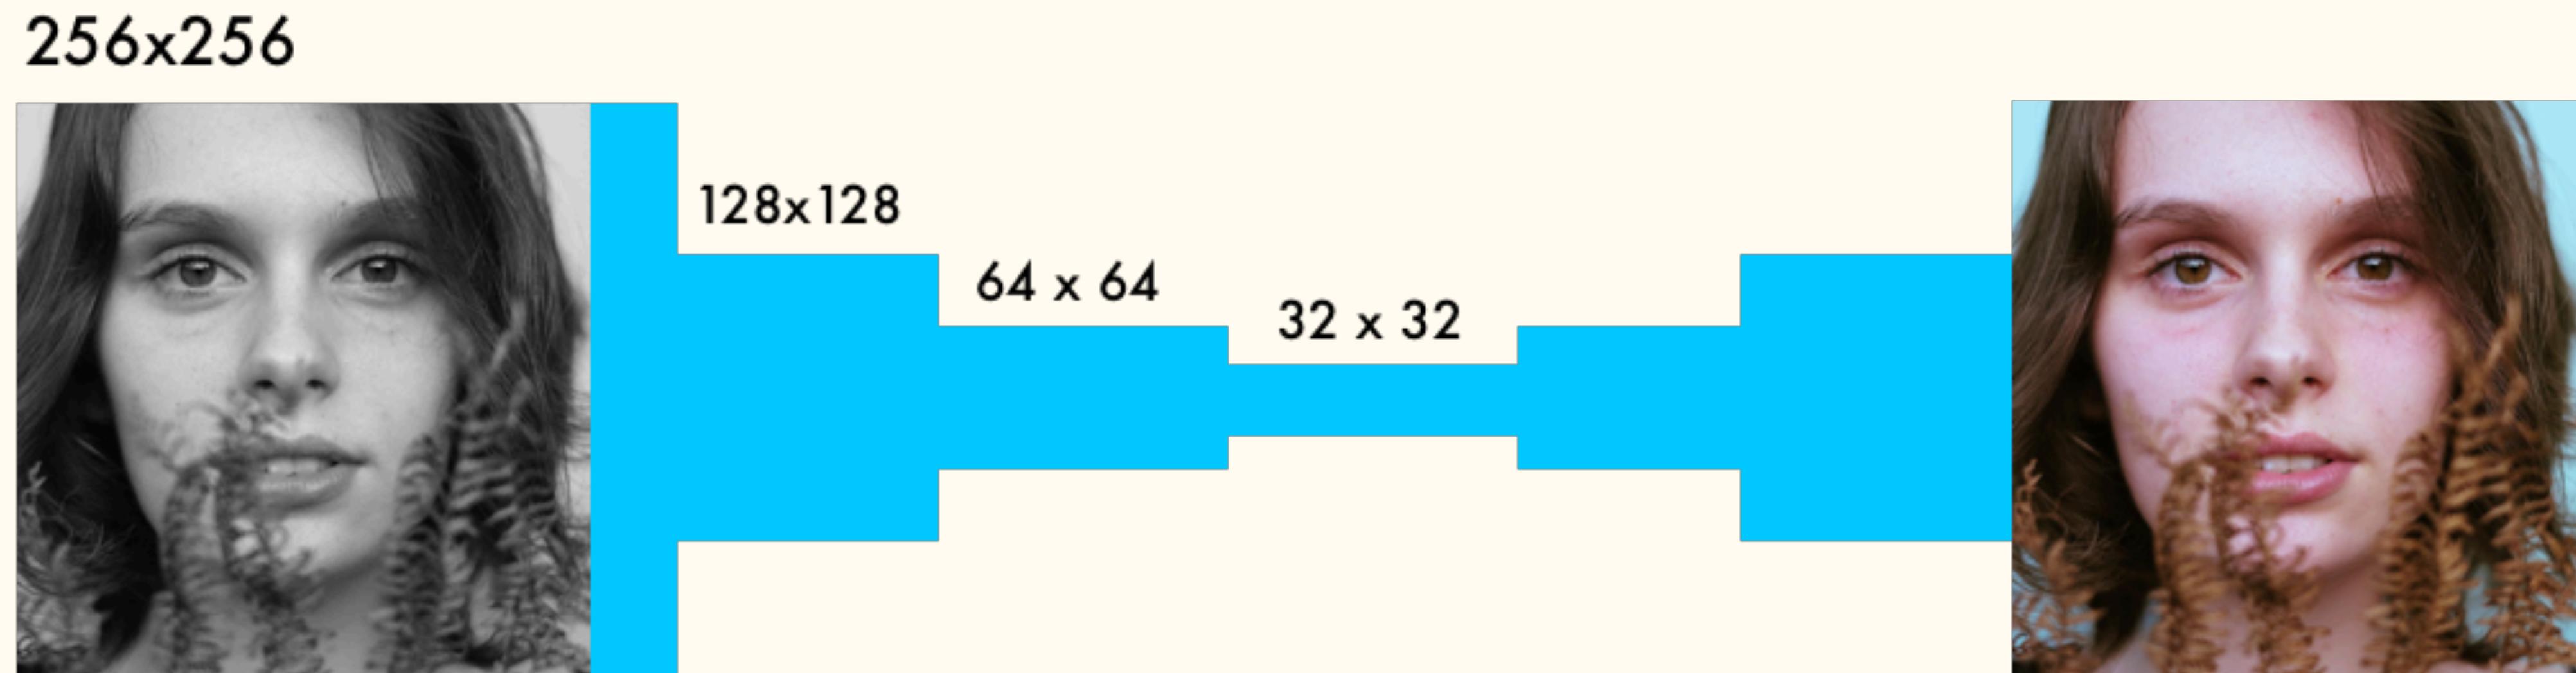


- The latent vector $z \sim (\mu, \sigma)$
- Compressed representation of input x
- $q_\Phi(z|x)$ encodes x to z
- $p_\theta(x|z)$ decodes z to x
- θ and Φ represent weights and biases
- Remember: we **sample** from the distribution $z \rightarrow$ lossy reconstruction

VAE Applications

Colorize B&W images

- Encoder-Decoder architecture
- Inception ResNet V2 Classifier



Generative Adversarial Network

- Works similarly to VAEs
 - ▶ Uses encoder-decoder system
 - ▶ Converts latent vectors into samples
 - ▶ ‘Generator’ instead of decoder
- But **trains** in a **different** way:
 - ▶ Passes random vectors into the generator
 - ▶ Directly evaluates the outputs on how well they follow the expected distribution

Generative Adversarial Network

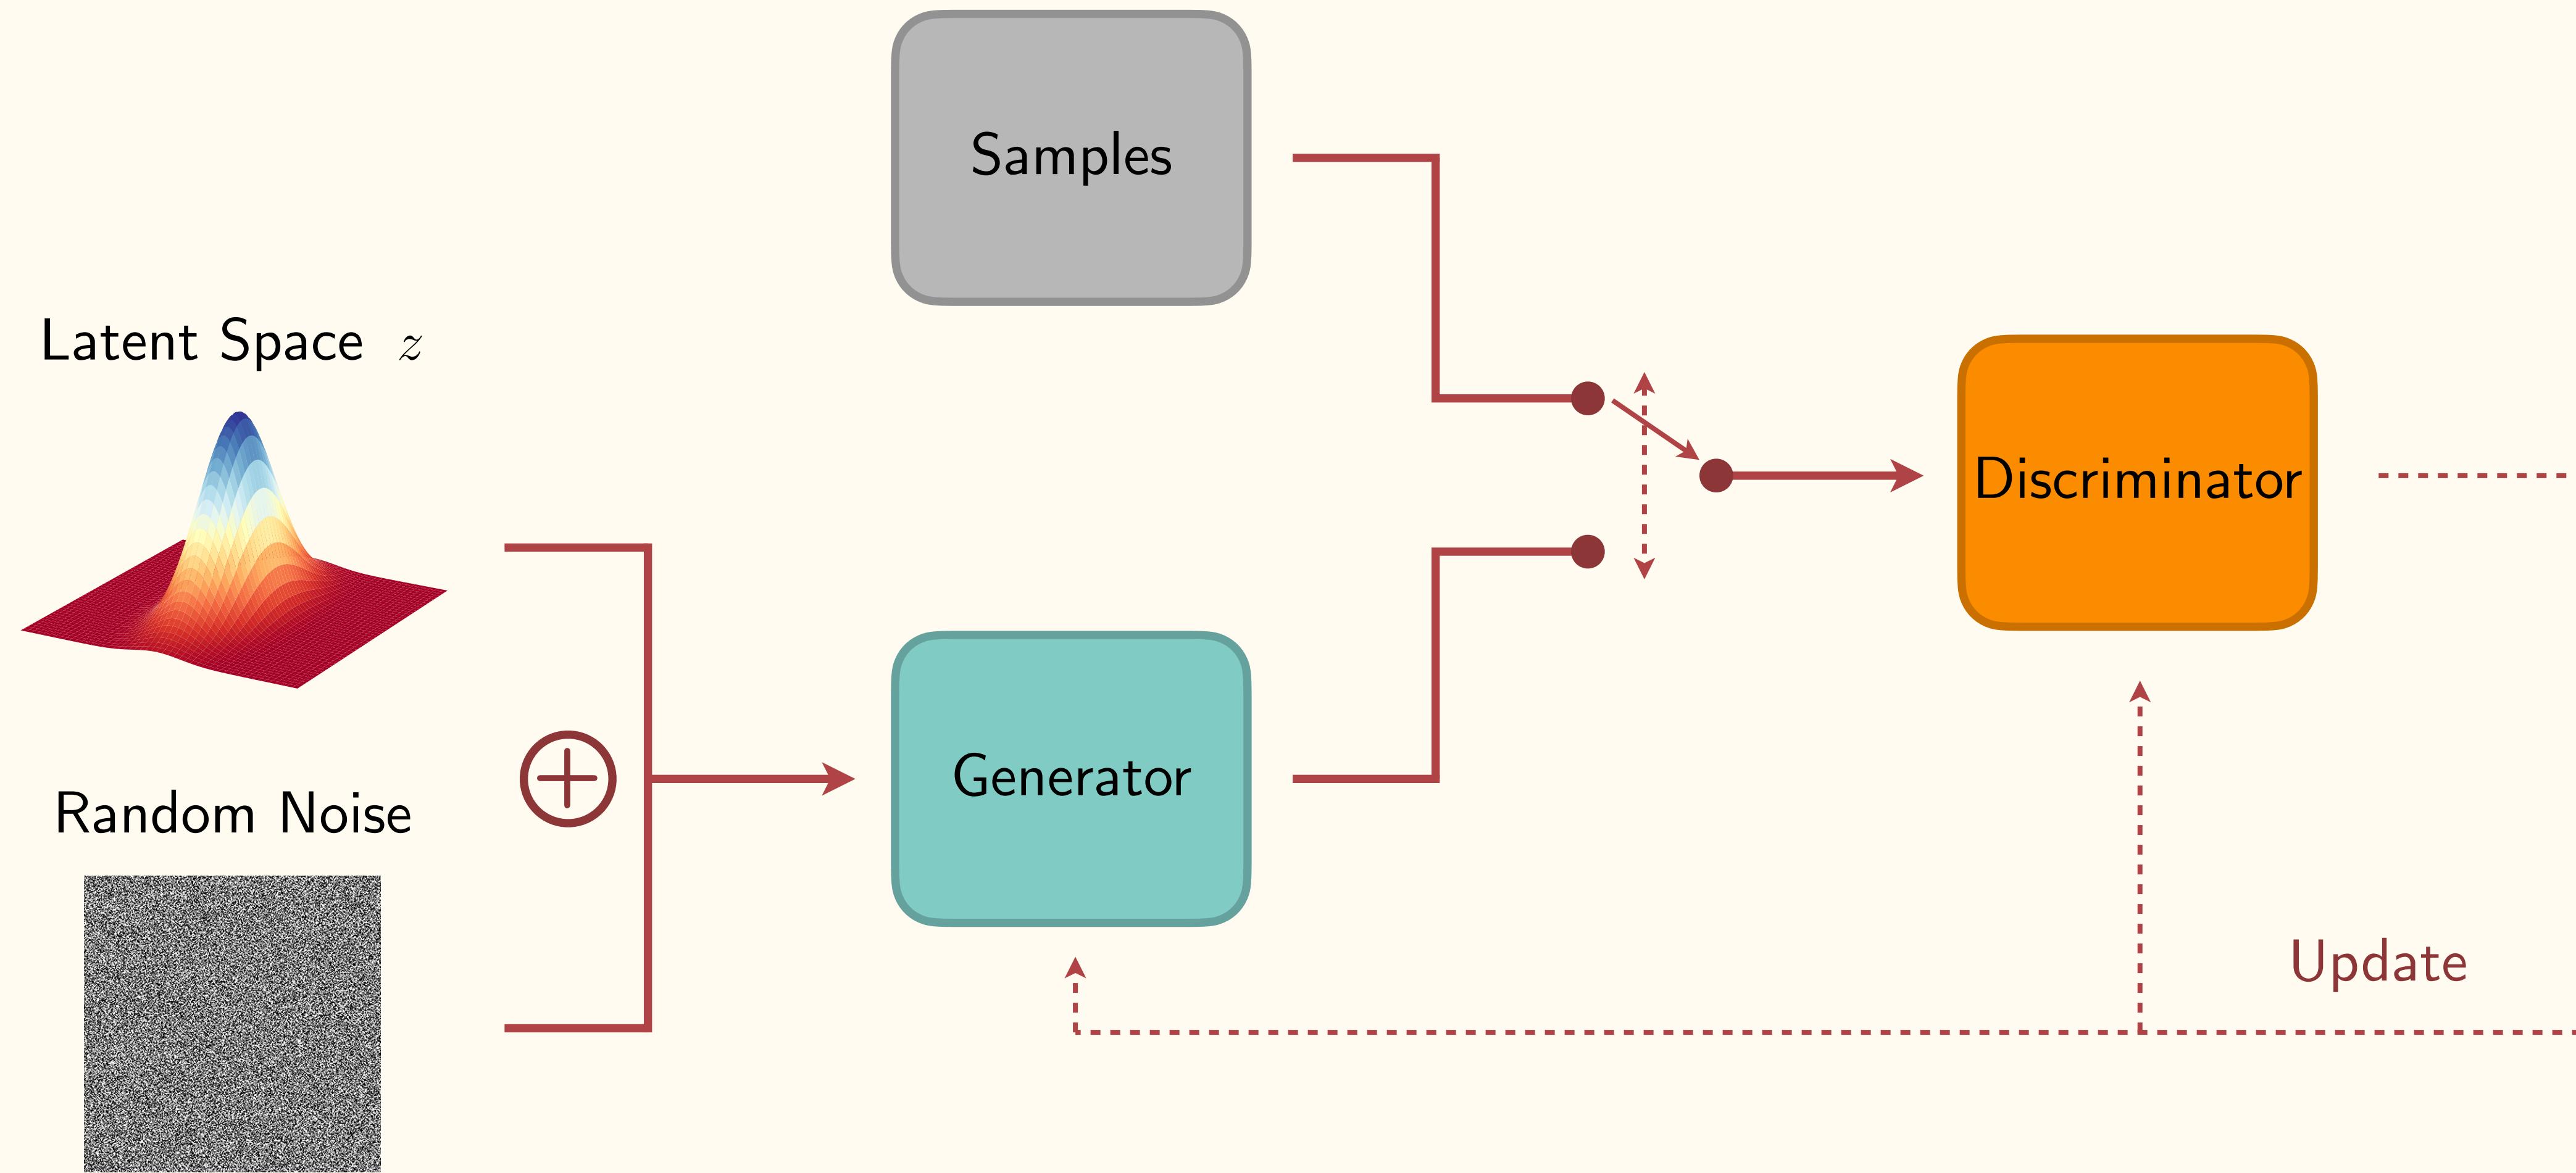
- Essentially create a loss function to measure:
 - How well generated samples match the training samples
 - Use that loss function to optimise the model
- How to make such a loss-function ?
 - You don't.
 - GAN learns the loss function from the data.



Generative Adversarial Network

- **Generator:** takes random vectors → generates synthetic samples
- **Discriminator:** tries to *distinguish* between 
 - The generated samples from real training samples
 - It takes a sample as input and outputs a **probability** of being a real sample
 - ▶ This acts as a loss function for the generator.

Generative Adversarial Network



Generative Adversarial Network

- **Generator** and **discriminator** are run *simultaneously*.
- Generator tries to ‘fool’ discriminator by presenting realistic samples
- Discriminator tries to get better at distinguishing real from fake samples
- Min-max game with value function $V(G, D)$:

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{data}(x)}[\log D(x)] + \mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z)))]$$

Generative *Adversarial* Networks

GAN Applications

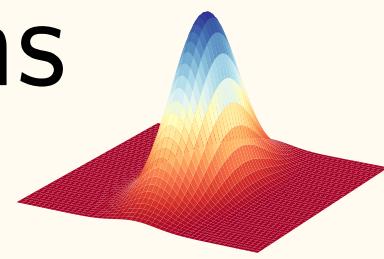
Colorize B&W images

- Even more advanced and realistic
- GAN based architecture



VAE vs GANs

Roughly:

- GANs tend to produce higher-quality samples 
- VAEs tend to produce higher-quality distributions 
- *Individual* samples generated by GANs more closely resemble training samples
- *Range* of samples generated by VAE more closely match range of training samples

Very active field of research → Things changing constantly !

Progression of GANs

- 2014: Original GAN
- 2015: Deep Convolutional GAN
- 2016: Coupled GAN
- 2017: Progressively Growing GAN
- 2018: Style-based GAN
- 2019: StyleGAN 2
- 2021: ?



2014



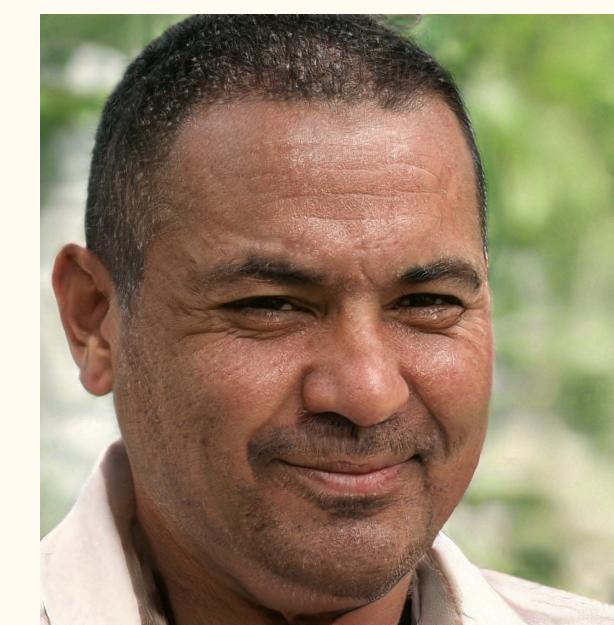
2015



2016



2017

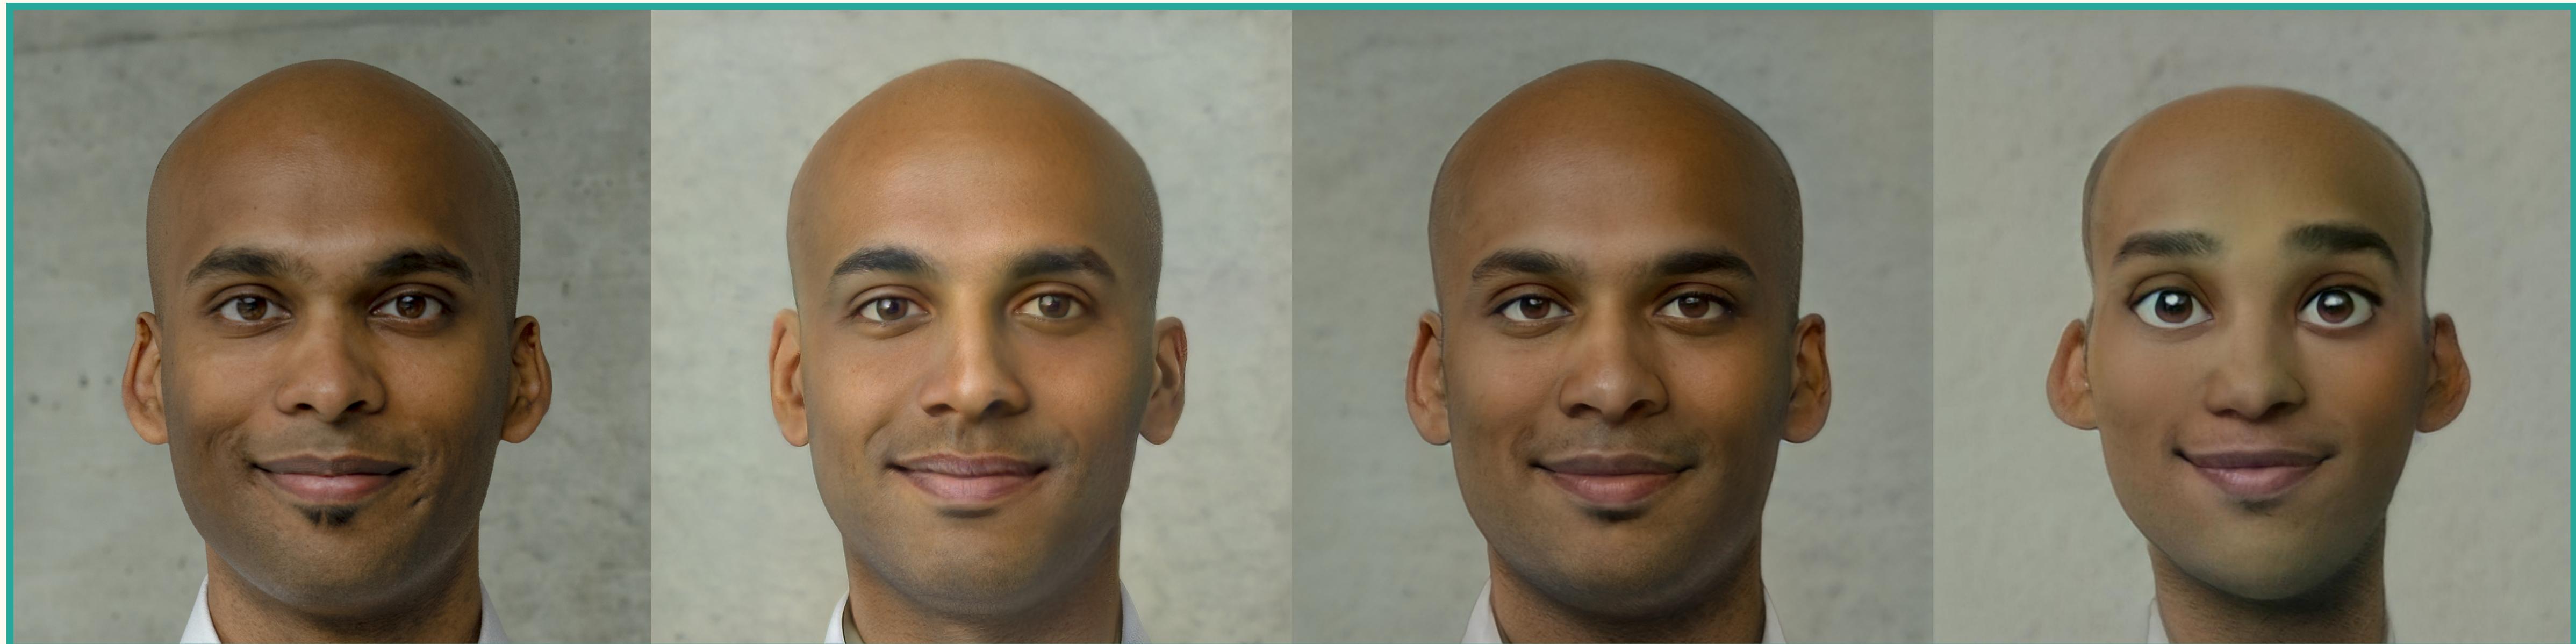


2018



2019

StyleGAN 2



Source input image

Projection into
FFHQ-StyleGAN2's
 W Latent Space

Projection into
FFHQ-StyleGAN2's
 $W+$ Latent Space

Projection into
Toon-StyleGAN2's
Latent Space

Thank you !



Room 207-2, Idiap Research Institute



www.idiap.ch/~esarkar/



+41 78 82 50 754



eklavya.sarkar@idiap.ch



Generative Adversarial Network

